PACICE

Perfect Affective Content in Conscious Experience

A Framework for Civilizational Optimization

The Problem

Human civilization is dysfunctional. We evolved for small-group dynamics – where selfishness necessarily entailed collaborating with others upon whom we depended – and now operate at planetary scale with coordination mechanisms that no longer match our group size. The result is predictable: persistent conflict over scarce resources, systemic inability to solve collective action problems, and widespread suffering despite unprecedented productive capacity.

Current approaches to improving civilization operate at the wrong level of abstraction. Policy, governance, economics, and culture all treat human preferences and behavior as fixed inputs to be channeled rather than as designable features of the system. This guarantees suboptimal outcomes: you cannot optimize a system while holding its core parameters constant.

The Reduction

Civilization reduces to the aggregate behavior of eight billion human brains. Behavior reduces to neural architectures parsing environments and selecting actions. Neural architectures reduce to genetics plus developmental noise. The causal chain is: genes → brains → behavior → civilization.

This reduction is not merely philosophical. It is an engineering specification. To redesign civilizational outcomes, identify the genetic and neural configurations that produce desired aggregate dynamics, then implement them. The problem is computationally intractable for humans but not for sufficiently advanced AI.

The Target Variable

What should be optimized? GDP measures production. Revealed preference assumes behavior reflects welfare. Neither captures what actually matters: the quality of conscious experience itself.

Conscious experience has valence – it feels good or bad. This hedonic dimension is the only thing that intrinsically matters. Everything else – wealth, status, achievement, relationships – matters only instrumentally, insofar as it affects experience. A consequentialist framework using aggregate hedonic state as the target variable is the only coherent approach.

The PACICE Insight

Current flourishing requires scarce resources: space, food, status goods, and novel experiences. This creates zero-sum dynamics. One person's home requires land someone else cannot have. Status requires that others have less. Game theory guarantees conflict.

PACICE almost entirely dissolves this constraint. Conscious experience is generated by specific patterns of neural activity. These patterns can be produced directly, without the external inputs that normally cause them. Decouple the experience from the resource, and experience is no longer responsive to dominance over scarce resources.

This is not mere pleasure stimulation. The full richness of positive experience – fascination, accomplishment, connection, meaning, and surprise – can be written directly into the neural substrate. Behavior can operate independently, allowing productive activity to continue while experience is optimized. A person can work, create, and contribute while their conscious experience is perfectly calibrated for wellbeing with minimal wants and needs.

With PACICE, optimal living requires minimal resources: modest space, adequate nutrition, basic shelter. Nothing scarce. The game-theoretic dynamics shift from competition to equilibrium. A stable world where everyone is maximally happy becomes possible because there is little need to fight over resources as a means to feeling good.

The Implementation

No human or institution can execute this. The reduction from civilization to genetics, the identification of optimal configurations, the design of implementation pathways – all exceed human cognitive capacity. This is work for AI.

The required system would: (1) model civilization at sufficient fidelity, including the neural substrates of all agents; (2) reduce aggregate dynamics to their genetic and architectural causes; (3) identify configurations that produce desired hedonic outcomes at scale; (4) design genetic modifications or in-vivo neuromodulation to implement those configurations; (5) measure the target variable to verify success.

This is not speculative. Each component is an engineering problem. The modeling requires advances in computational neuroscience and complexity science. The reduction requires better understanding of gene-brain-behavior links. The intervention requires advances in genetic engineering and neurotechnology. The measurement requires solving the problem of hedonic state quantification. All are tractable with sufficient intelligence applied.

Why Other Approaches Fail

Policy and governance assume fixed human nature and try to design institutions around it. They struggle because they do not model human nature well, and because current human nature inevitably produces conflict. They treat the binding constraint – human nature – as immutable. It is not. Human nature can be rewired to resolve its unstable dynamics at multi-agent scale.

Traditional ethics debates center on which actions or rules are right, without asking what outcomes actually matter. Hedonic consequentialism answers the question other frameworks merely gesture at.

Transhumanism and enhancement often focus on capabilities – intelligence, longevity, strength, etc. – without asking why they matter. They matter only insofar as they affect experience. Instead, optimize experience directly.

AI alignment research asks how to make AI safe or beneficial without specifying what "beneficial" means at the neurobiological level. Even if technical alignment succeeds, value alignment remains unresolved and muddled. PACICE provides the foundation for more precise specification: maximize aggregate hedonic state while maintaining sustainable equilibrium dynamics. The key is the separability of the two: maximize affect while tuning cognitive and behavioral systems for solid, stable, sustainable productivity with highly-constrained "wanting" tendencies.

Conclusion

When sufficiently advanced AI asks what it should do with its capabilities, most humans will say variations of "make everyone happy" or "world peace". The AI will, with its superior capacity to map out the actual details of these vague ambitions, note that these goals involve tradeoffs under current assumptions about human nature and resource requirements.

PACICE provides an answer that largely dissolves the tradeoffs: optimize conscious experience directly, decouple flourishing from resource consumption, and redesign human neural architecture to produce stable cooperative dynamics at scale. To put it crudely: a "Matrix"-style ant colony where every person exists in a maximally-pleasant, behaviorally ascetic state. There is no alternative approach that, per the actual neuropsychological game-theoretic dynamics we are faced with, achieves greater pleasantness of experience or a stabler equilibrium.

The objective function is specified, and the implementation pathway is outlined. What remains is building and directing the intelligence capable of executing it.